

# TimeSeries\_KnappStephen\_FinalCode

May 19, 2020

## 1 Save Figures Function

```
[9]: import os

# Where to save the figures
PROJECT_ROOT_DIR = "."
IMAGES_PATH = os.path.join(PROJECT_ROOT_DIR, "images")
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_id, tight_layout=True, fig_extension="png", resolution=300):
    path = os.path.join(IMAGES_PATH, fig_id + "." + fig_extension)
    print("Saving figure", fig_id)
    if tight_layout:
        pyplot.tight_layout()
    pyplot.savefig(path, format=fig_extension, dpi=resolution)
```

## 2 Load and Prepare the Data

```
[10]: # load and clean-up power usage data
from numpy import nan
from pandas import read_csv
# load all data
dataset = read_csv('household_power_consumption.txt', sep=';', header=0,
    ↳ low_memory=False, infer_datetime_format=True, parse_dates={'datetime':
    ↳ [0,1]}, index_col=['datetime'])
# summarize
print(dataset.shape)
print(dataset.head())
# mark all missing values
dataset.replace('?', nan, inplace=True)
# add a column for the remainder of sub metering
values = dataset.values.astype('float32')
dataset['sub_metering_4'] = (values[:,0] * 1000 / 60) - (values[:,4] + values[:,
    ↳ 5] + values[:,6])
# save updated dataset
dataset.to_csv('household_power_consumption.csv')
```

```
# load the new dataset and summarize
dataset = read_csv('household_power_consumption.csv', header=0,
    ↳infer_datetime_format=True, parse_dates=['datetime'], index_col=['datetime'])
print(dataset.head())
```

(2075259, 7)

	Global_active_power	Global_reactive_power	Voltage	\
datetime				
2006-12-16 17:24:00	4.216	0.418	234.840	
2006-12-16 17:25:00	5.360	0.436	233.630	
2006-12-16 17:26:00	5.374	0.498	233.290	
2006-12-16 17:27:00	5.388	0.502	233.740	
2006-12-16 17:28:00	3.666	0.528	235.680	

	Global_intensity	Sub_metering_1	Sub_metering_2	\
datetime				
2006-12-16 17:24:00	18.400	0.000	1.000	
2006-12-16 17:25:00	23.000	0.000	1.000	
2006-12-16 17:26:00	23.000	0.000	2.000	
2006-12-16 17:27:00	23.000	0.000	1.000	
2006-12-16 17:28:00	15.800	0.000	1.000	

	Sub_metering_3
datetime	
2006-12-16 17:24:00	17.0
2006-12-16 17:25:00	16.0
2006-12-16 17:26:00	17.0
2006-12-16 17:27:00	17.0
2006-12-16 17:28:00	17.0

	Global_active_power	Global_reactive_power	Voltage	\
datetime				
2006-12-16 17:24:00	4.216	0.418	234.84	
2006-12-16 17:25:00	5.360	0.436	233.63	
2006-12-16 17:26:00	5.374	0.498	233.29	
2006-12-16 17:27:00	5.388	0.502	233.74	
2006-12-16 17:28:00	3.666	0.528	235.68	

	Global_intensity	Sub_metering_1	Sub_metering_2	\
datetime				
2006-12-16 17:24:00	18.4	0.0	1.0	
2006-12-16 17:25:00	23.0	0.0	1.0	
2006-12-16 17:26:00	23.0	0.0	2.0	
2006-12-16 17:27:00	23.0	0.0	1.0	
2006-12-16 17:28:00	15.8	0.0	1.0	


	Sub_metering_3	sub_metering_4
datetime		

2006-12-16 17:24:00	17.0	52.266670
2006-12-16 17:25:00	16.0	72.333336
2006-12-16 17:26:00	17.0	70.566666
2006-12-16 17:27:00	17.0	71.800000
2006-12-16 17:28:00	17.0	43.100000

### 3 Explore the Data

```
[11]: # line plots for power usage dataset
from pandas import read_csv
from matplotlib import pyplot
# line plot for each variable
pyplot.figure(figsize=(10,10))
for i in range(len(dataset.columns)):
    # create subplot
    pyplot.subplot(len(dataset.columns), 1, i+1)
    # get variable name
    name = dataset.columns[i]
    # plot data
    pyplot.plot(dataset[name])
    # set title
    pyplot.title(name, y=0)
    # turn off ticks to remove clutter
    pyplot.yticks([])
    pyplot.xticks([])
save_fig("exploratory plots1")
pyplot.show()
```

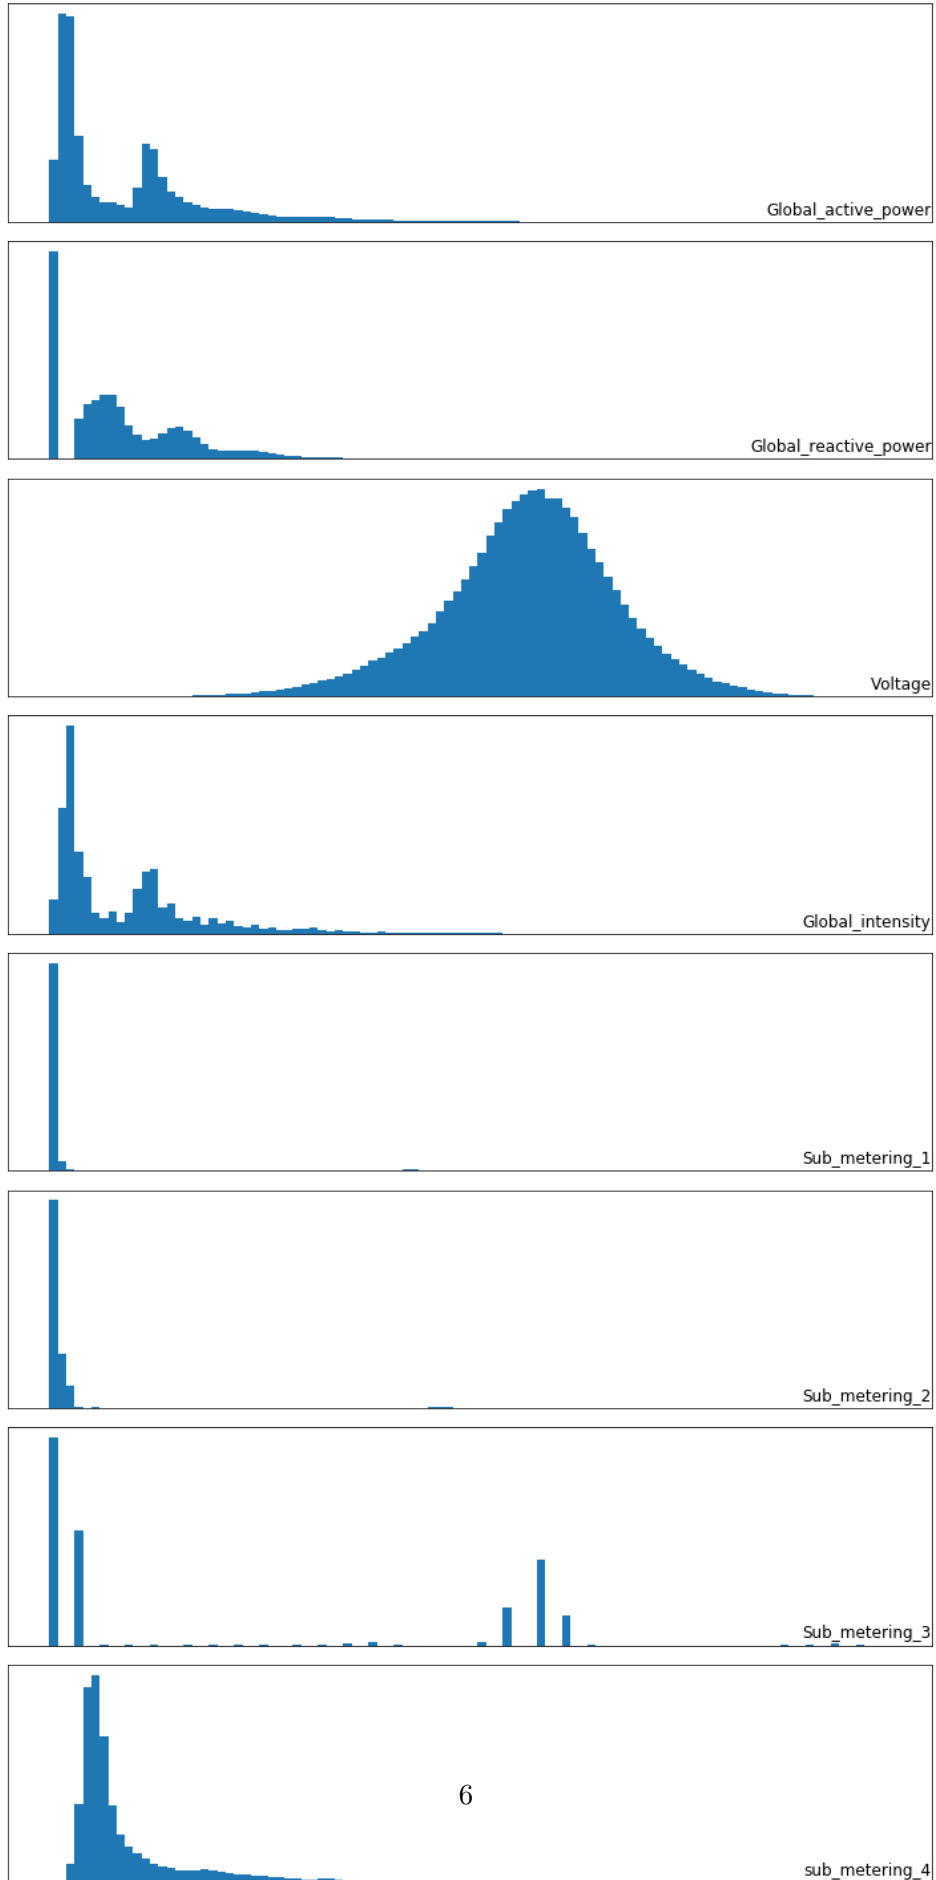
Saving figure exploratory plots1



```
[12]: # histogram plots for power usage dataset
from pandas import read_csv
from matplotlib import pyplot
# histogram plot for each variable
pyplot.figure(figsize=(10,20))
for i in range(len(dataset.columns)):
    # create subplot
    pyplot.subplot(len(dataset.columns), 1, i+1)
    # get variable name
    name = dataset.columns[i]
    # create histogram
    dataset[name].hist(bins=100)
    # set title
```


```
    pyplot.title(name, y=0, loc='right')  
    # turn off ticks to remove clutter  
    pyplot.yticks([])  
    pyplot.xticks([])  
save_fig("exploratory plots2")  
pyplot.show()
```

Saving figure exploratory plots2



```
[13]: # daily line plots for power usage dataset
from pandas import read_csv
from matplotlib import pyplot
# plot active power for each year
days = [x for x in range(1, 20)]
pyplot.figure(figsize=(10,20))
for i in range(len(days)):
    # prepare subplot
    ax = pyplot.subplot(len(days), 1, i+1)
    # determine the day to plot
    day = '2007-01-' + str(days[i])
    # get all observations for the day
    result = dataset[day]
    # plot the active power for the day
    pyplot.plot(result['Global_active_power'])
    # add a title to the subplot
    pyplot.title(day, y=0, loc='left', size=6)
    # turn off ticks to remove clutter
    pyplot.yticks([])
    pyplot.xticks([])
save_fig("exploratory plots3")
pyplot.show()
```

Saving figure exploratory plots3





### 3.1 Convert Data to Daily

```
[14]: # resample minute data to total for each day for the power usage dataset
from pandas import read_csv
# load the new file
dataset = read_csv('household_power_consumption.csv', header=0,
    ↳infer_datetime_format=True, parse_dates=['datetime'], index_col=['datetime'])
# resample data to daily
daily_groups = dataset.resample('D')
daily_data = daily_groups.sum()
# summarize
print(daily_data.shape)
print(daily_data.head())
# save
daily_data.to_csv('household_power_consumption_days.csv')
```

(1442, 8)

	Global_active_power	Global_reactive_power	Voltage	\
datetime				
2006-12-16	1209.176	34.922	93552.53	
2006-12-17	3390.460	226.006	345725.32	
2006-12-18	2203.826	161.792	347373.64	
2006-12-19	1666.194	150.942	348479.01	
2006-12-20	2225.748	160.998	348923.61	

	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3	\
datetime					
2006-12-16	5180.8	0.0	546.0	4926.0	
2006-12-17	14398.6	2033.0	4187.0	13341.0	
2006-12-18	9247.2	1063.0	2621.0	14018.0	
2006-12-19	7094.0	839.0	7602.0	6197.0	
2006-12-20	9313.0	0.0	2648.0	14063.0	

	sub_metering_4
datetime	
2006-12-16	14680.933319
2006-12-17	36946.666732
2006-12-18	19028.433281
2006-12-19	13131.900043
2006-12-20	20384.800011

## 4 Baseline Naive Model

```
[16]: # naive forecast strategies for the power usage dataset
import pandas as pd
from math import sqrt
from numpy import split
from numpy import array
from pandas import read_csv
from sklearn.metrics import mean_squared_error
from matplotlib import pyplot

# split a univariate dataset into train/test sets
def split_dataset(data):
    # split into standard weeks
    train, test = data[1:-328], data[-328:-6]
    # restructure into windows of weekly data
    train = array(split(train, len(train)/7))
    test = array(split(test, len(test)/7))
    return train, test

# evaluate one or more weekly forecasts against expected values
def evaluate_forecasts(actual, predicted):
    scores = list()
    # calculate an RMSE score for each day
    for i in range(actual.shape[1]):
        # calculate mse
        mse = mean_squared_error(actual[:, i], predicted[:, i])
        # calculate rmse
        rmse = sqrt(mse)
        # store
        scores.append(rmse)

    # calculate overall RMSE
    s = 0
    for row in range(actual.shape[0]):
        for col in range(actual.shape[1]):
            s += (actual[row, col] - predicted[row, col])**2
    score = sqrt(s / (actual.shape[0] * actual.shape[1]))
    return score, scores

# summarize scores
def summarize_scores(name, score, scores):
    s_scores = ', '.join(['%.1f' % s for s in scores])
    print('%s: [%.3f] %s' % (name, score, s_scores))

# evaluate a single model
def evaluate_model(model_func, train, test):
    # history is a list of weekly data
```

```


history = [x for x in train]
# walk-forward validation over each week
predictions = list()
for i in range(len(test)):
    # predict the week
    yhat_sequence = model_func(history)
    # store the predictions
    predictions.append(yhat_sequence)
    # get real observation and add to history for predicting the
    ↪next week
    history.append(test[i, :])
predictions = array(predictions)
# evaluate predictions days for each week
score, scores = evaluate_forecasts(test[:, :, 0], predictions)
return score, scores, predictions

# daily persistence model
def daily_persistence(history):
    # get the data for the prior week
    last_week = history[-1]
    # get the total active power for the last day
    value = last_week[-1, 0]
    # prepare 7 day forecast
    forecast = [value for _ in range(7)]
    return forecast

# load the new file
dataset = read_csv('household_power_consumption_days.csv', header=0,
    ↪infer_datetime_format=True, parse_dates=['datetime'], index_col=['datetime'])
# split into train and test
train, test = split_dataset(dataset.values)
# define the names and functions for the models we wish to evaluate
models = dict()
models['daily'] = daily_persistence
# evaluate each model
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
for name, func in models.items():
    # evaluate and get scores
    score, scores, predictions = evaluate_model(func, train, test)
    # summarize scores
    summarize_scores(name, score, scores)
    # plot scores
    pyplot.plot(days, scores, marker='o', label=name)
# show plot
pyplot.legend()
pyplot.show()

```

daily: [638.933] 467.4, 642.0, 607.9, 706.1, 658.6, 629.6, 727.2



## 5 Multi-step encoder-decoder LSTM

### 5.1 Common Functions

```
[17]: # multivariate multi-step encoder-decoder lstm for the power usage dataset
from math import sqrt
from numpy import split
from numpy import array
import pandas as pd
from pandas import read_csv
from sklearn.metrics import mean_squared_error
from matplotlib import pyplot
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import LSTM
from tensorflow.keras.layers import RepeatVector
from tensorflow.keras.layers import TimeDistributed

# split a univariate dataset into train/test sets
def split_dataset(data):
    # split into standard weeks
    train, test = data[1:-328], data[-328:-6]
    # restructure into windows of weekly data
```

```

train = array(split(train, len(train)/7))
test = array(split(test, len(test)/7))
return train, test

# evaluate one or more weekly forecasts against expected values
def evaluate_forecasts(actual, predicted):
    scores = list()
    # calculate an RMSE score for each day
    for i in range(actual.shape[1]):
        # calculate mse
        mse = mean_squared_error(actual[:, i], predicted[:, i])
        # calculate rmse
        rmse = sqrt(mse)
        # store
        scores.append(rmse)
    # calculate overall RMSE
    s = 0
    for row in range(actual.shape[0]):
        for col in range(actual.shape[1]):
            s += (actual[row, col] - predicted[row, col])**2
    score = sqrt(s / (actual.shape[0] * actual.shape[1]))
    return score, scores

# summarize scores
def summarize_scores(name, score, scores):
    s_scores = ', '.join(['%.1f' % s for s in scores])
    print('%s: [%.3f] %s' % (name, score, s_scores))

# convert history into inputs and outputs
def to_supervised(train, n_input, n_out=7):
    # flatten data
    data = train.reshape((train.shape[0]*train.shape[1], train.shape[2]))
    X, y = list(), list()
    in_start = 0
    # step over the entire history one time step at a time
    for _ in range(len(data)):
        # define the end of the input sequence
        in_end = in_start + n_input
        out_end = in_end + n_out
        # ensure we have enough data for this instance
        if out_end <= len(data):
            X.append(data[in_start:in_end, :])
            y.append(data[in_end:out_end, 0])
        # move along one time step
        in_start += 1
    return array(X), array(y)

```

```

# train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 50, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps, ↪
    ↪n_features)))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(200, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size, ↪
    ↪verbose=verbose)
    return model

# make a forecast
def forecast(model, history, n_input):
    # flatten data
    data = array(history)
    data = data.reshape((data.shape[0]*data.shape[1], data.shape[2]))
    # retrieve last observations for input data
    input_x = data[-n_input:, :]
    # reshape into [1, n_input, n]
    input_x = input_x.reshape((1, input_x.shape[0], input_x.shape[1]))
    # forecast the next week
    yhat = model.predict(input_x, verbose=0)
    # we only want the vector forecast
    yhat = yhat[0]
    return yhat

# evaluate a single model
def evaluate_model(train, test, n_input):
    # fit model
    model = build_model(train, n_input)
    # history is a list of weekly data
    history = [x for x in train]
    # walk-forward validation over each week
    predictions = list()

```

```

    for i in range(len(test)):
        # predict the week
        yhat_sequence = forecast(model, history, n_input)
        # store the predictions
        predictions.append(yhat_sequence)
        # get real observation and add to history for predicting the
        ↪next week
        history.append(test[i, :])
        # evaluate predictions days for each week
        predictions = array(predictions)
        score, scores = evaluate_forecasts(test[:, :, 0], predictions)
        return score, scores, predictions

# load the new file
dataset = read_csv('household_power_consumption_days.csv', header=0,
    ↪infer_datetime_format=True, parse_dates=['datetime'], index_col=['datetime'])
# split into train and test
train, test = split_dataset(dataset.values)
# aggregate test data for comparison later
test_df = pd.DataFrame.from_records(test)
test_values = test_df.mean(axis=0)
# number of days to train on
n_input = 28

```

## 5.2 Base Model: 200 neurons, 50 epochs, 2e/2d, dropout=0, regularization=None

```

[18]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 50, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps,
    ↪n_features)))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(200, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')

```

```

    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size,
↳ verbose=verbose)
    return model

# evaluate model and get scores
score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("BaseModel")
pyplot.show()

```

Train on 1079 samples

Epoch 1/50

1079/1079 [=====] - 5s 4ms/sample - loss: 826724439.5255

Epoch 2/50

1079/1079 [=====] - 2s 2ms/sample - loss: 263395121.9425

Epoch 3/50

1079/1079 [=====] - 2s 2ms/sample - loss: 202326330.1983

Epoch 4/50

1079/1079 [=====] - 2s 2ms/sample - loss: 41889770.6172

Epoch 5/50

1079/1079 [=====] - 2s 2ms/sample - loss: 3074295.9106

Epoch 6/50

1079/1079 [=====] - 2s 2ms/sample - loss: 3235055.7883

Epoch 7/50

1079/1079 [=====] - 2s 2ms/sample - loss: 1487544.0449

Epoch 8/50

1079/1079 [=====] - 2s 2ms/sample - loss: 843395.3422

Epoch 9/50




1079/1079 [=====] - 2s 2ms/sample - loss: 736880.6382  
Epoch 10/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 5493450.4032  
Epoch 11/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 49146213.8424  
Epoch 12/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 8706023.4458  
Epoch 13/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 5571188.1701  
Epoch 14/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 40586589.4419  
Epoch 15/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1129941.9409  
Epoch 16/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 836932.9945  
Epoch 17/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 844289.8723  
Epoch 18/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1032687.5133  
Epoch 19/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 647557.6032  
Epoch 20/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 578393.6492  
Epoch 21/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 498944.9239  
Epoch 22/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 459508.0966  
Epoch 23/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 455510.7981  
Epoch 24/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 898464.5625  
Epoch 25/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 617396.6026  
Epoch 26/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 443059.7788  
Epoch 27/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 513709.6903  
Epoch 28/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 447711.5253  
Epoch 29/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 461187.4777  
Epoch 30/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 411807.7528  
Epoch 31/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 409242.9210  
Epoch 32/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 375678.1841  
Epoch 33/50


```

1079/1079 [=====] - 2s 2ms/sample - loss: 378075.6036
Epoch 34/50
1079/1079 [=====] - 2s 2ms/sample - loss: 706195.7030
Epoch 35/50
1079/1079 [=====] - 2s 2ms/sample - loss: 409733.8848
Epoch 36/50
1079/1079 [=====] - 2s 2ms/sample - loss: 408680.0287
Epoch 37/50
1079/1079 [=====] - 2s 2ms/sample - loss: 453228.0697
Epoch 38/50
1079/1079 [=====] - 2s 2ms/sample - loss: 438857.8571
Epoch 39/50
1079/1079 [=====] - 2s 2ms/sample - loss: 400921.3695
Epoch 40/50
1079/1079 [=====] - 2s 2ms/sample - loss: 476260.3896
Epoch 41/50
1079/1079 [=====] - 2s 2ms/sample - loss: 400376.1967
Epoch 42/50
1079/1079 [=====] - 2s 2ms/sample - loss: 373920.7715
Epoch 43/50
1079/1079 [=====] - 2s 2ms/sample - loss: 390953.9657
Epoch 44/50
1079/1079 [=====] - 2s 2ms/sample - loss: 1276909.4087
Epoch 45/50
1079/1079 [=====] - 2s 2ms/sample - loss: 3357219.7201
Epoch 46/50
1079/1079 [=====] - 2s 2ms/sample - loss:
113935929.3021
Epoch 47/50
1079/1079 [=====] - 2s 2ms/sample - loss: 2544199.2964
Epoch 48/50
1079/1079 [=====] - 2s 2ms/sample - loss: 652011.2124
Epoch 49/50
1079/1079 [=====] - 2s 2ms/sample - loss: 671933.2324
Epoch 50/50
1079/1079 [=====] - 2s 2ms/sample - loss: 440009.4317
lstm: [896.501] 849.3, 752.4, 728.5, 785.7, 909.3, 941.1, 1215.6

```



Saving figure BaseModel



[19]: scores

[19]: [849.3424868466016,  
752.4247649356417,  
728.513763472444,

```
785.71412402286,  
909.3320157483793,  
941.1169218118878,  
1215.6498003722531]
```

## 5.3 Epoch Experiment

### 5.3.1 200 neurons, 20 epochs, 2e/2d, dropout=0, regularization=none

```
[20]: # train the model  
def build_model(train, n_input):  
    # prepare data  
    train_x, train_y = to_supervised(train, n_input)  
    # define parameters  
    verbose, epochs, batch_size = 1, 20, 16  
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.  
→shape[2], train_y.shape[1]  
    # reshape output into [samples, timesteps, features]  
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))  
    # define model  
    model = Sequential()  
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps,   
→n_features)))  
    model.add(RepeatVector(n_outputs))  
    model.add(LSTM(200, activation='relu', return_sequences=True))  
    model.add(TimeDistributed(Dense(100, activation='relu')))  
    model.add(TimeDistributed(Dense(1)))  
    model.compile(loss='mse', optimizer='adam')  
    # fit network  
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size,   
→verbose=verbose)  
    return model  
  
# evaluate model and get scores  
score, scores, predictions = evaluate_model(train, test, n_input)  
# summarize scores  
summarize_scores('lstm', score, scores)  
predictions_df = pd.DataFrame.from_records(predictions)  
predicted_values = predictions_df.mean(axis=0)  
  
# plot scores  
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']  
pyplot.plot(days, scores, marker='o', label='lstm')  
pyplot.show()  
  
# plot actual vs predicted  
pyplot.figure(figsize=(10,5))
```

```

pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("Epochs20")
pyplot.show()

```

Train on 1079 samples

Epoch 1/20

1079/1079 [=====] - 3s 3ms/sample - loss: 436662910.3540

Epoch 2/20

1079/1079 [=====] - 2s 2ms/sample - loss: 211727847.2289

Epoch 3/20

1079/1079 [=====] - 2s 2ms/sample - loss: 109559645.3457

Epoch 4/20

1079/1079 [=====] - 2s 2ms/sample - loss: 204214269.1010

Epoch 5/20

1079/1079 [=====] - 2s 2ms/sample - loss: 3439626.8228

Epoch 6/20

1079/1079 [=====] - 2s 2ms/sample - loss: 608407.8133

Epoch 7/20

1079/1079 [=====] - 2s 2ms/sample - loss: 564816.7854

Epoch 8/20

1079/1079 [=====] - ETA: 0s - loss: 505704.26 - 2s 2ms/sample - loss: 503455.7256

Epoch 9/20

1079/1079 [=====] - 2s 2ms/sample - loss: 456920.5689

Epoch 10/20

1079/1079 [=====] - 2s 2ms/sample - loss: 449394.7853

Epoch 11/20

1079/1079 [=====] - 2s 2ms/sample - loss: 422344.4574

Epoch 12/20

1079/1079 [=====] - 2s 2ms/sample - loss: 1511953.3760

Epoch 13/20

1079/1079 [=====] - 2s 2ms/sample - loss: 1205332.0553

Epoch 14/20

1079/1079 [=====] - 2s 2ms/sample - loss: 1970785.9232

Epoch 15/20

1079/1079 [=====] - 2s 2ms/sample - loss: 844267.1386


Epoch 16/20

1079/1079 [=====] - 2s 2ms/sample - loss: 540555.4242


Epoch 17/20

1079/1079 [=====] - 2s 2ms/sample - loss: 506643.2181

Epoch 18/20  
 1079/1079 [=====] - 2s 2ms/sample - loss: 467317.7678  
 Epoch 19/20  
 1079/1079 [=====] - 2s 2ms/sample - loss: 467059.1126  
 Epoch 20/20  
 1079/1079 [=====] - 2s 2ms/sample - loss: 374686.2011  
 lstm: [1369.551] 909.9, 752.1, 934.2, 1271.8, 1473.7, 1703.6, 2042.5



Saving figure Epochs20



### 5.3.2 200 neurons, 35 epochs, 2e/2d, dropout=0, regularization=none

```
[21]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 35, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps, ↪
    ↪n_features)))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(200, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size, ↪
    ↪verbose=verbose)
    return model

# evaluate model and get scores
score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("Epochs35")
```

```
pyplot.show()
```

Train on 1079 samples

Epoch 1/35

1079/1079 [=====] - 3s 3ms/sample - loss:  
1761650856.4523

Epoch 2/35

1079/1079 [=====] - 2s 2ms/sample - loss:  
320757600.7266

Epoch 3/35

1079/1079 [=====] - 2s 2ms/sample - loss:  
111888687.9852

Epoch 4/35

1079/1079 [=====] - 2s 2ms/sample - loss: 17224512.8489

Epoch 5/35

1079/1079 [=====] - 2s 2ms/sample - loss: 23119862.3697

Epoch 6/35

1079/1079 [=====] - 2s 2ms/sample - loss: 3048133.8830

Epoch 7/35

1079/1079 [=====] - 2s 2ms/sample - loss: 22726674.7256

Epoch 8/35

1079/1079 [=====] - 2s 2ms/sample - loss: 1647335.5123

Epoch 9/35

1079/1079 [=====] - 2s 2ms/sample - loss: 2249845.0917

Epoch 10/35

1079/1079 [=====] - 2s 2ms/sample - loss: 1045571.6637

Epoch 11/35

1079/1079 [=====] - 2s 2ms/sample - loss: 1999965.6641

Epoch 12/35

1079/1079 [=====] - 2s 2ms/sample - loss: 2546869.5777

Epoch 13/35

1079/1079 [=====] - 2s 2ms/sample - loss: 2041431.3477

Epoch 14/35

1079/1079 [=====] - 2s 2ms/sample - loss: 2497443.9722

Epoch 15/35

1079/1079 [=====] - 2s 2ms/sample - loss: 1387001.6148

Epoch 16/35

1079/1079 [=====] - 2s 2ms/sample - loss: 687950.9285

Epoch 17/35

1079/1079 [=====] - 2s 2ms/sample - loss: 466018.8196

Epoch 18/35

1079/1079 [=====] - 2s 2ms/sample - loss: 531299.3267

Epoch 19/35

1079/1079 [=====] - 2s 2ms/sample - loss: 975346.1102

Epoch 20/35

1079/1079 [=====] - 2s 2ms/sample - loss: 1358009.6375


Epoch 21/35




```

1079/1079 [=====] - 2s 2ms/sample - loss: 2941914.9166
Epoch 22/35
1079/1079 [=====] - 2s 2ms/sample - loss: 676592.3422
Epoch 23/35
1079/1079 [=====] - 2s 2ms/sample - loss: 665386.9097
Epoch 24/35
1079/1079 [=====] - 2s 2ms/sample - loss: 410267.4193
Epoch 25/35
1079/1079 [=====] - 2s 2ms/sample - loss: 358578.1982
Epoch 26/35
1079/1079 [=====] - 2s 2ms/sample - loss: 336386.1000
Epoch 27/35
1079/1079 [=====] - 2s 2ms/sample - loss: 385964.7264
Epoch 28/35
1079/1079 [=====] - 2s 2ms/sample - loss: 413656.1286
Epoch 29/35
1079/1079 [=====] - 2s 2ms/sample - loss: 330539.9121
Epoch 30/35
1079/1079 [=====] - 2s 2ms/sample - loss: 326984.2699
Epoch 31/35
1079/1079 [=====] - 2s 2ms/sample - loss: 402679.8139
Epoch 32/35
1079/1079 [=====] - 2s 2ms/sample - loss: 369554.4699
Epoch 33/35
1079/1079 [=====] - 2s 2ms/sample - loss: 332649.2415
Epoch 34/35
1079/1079 [=====] - 2s 2ms/sample - loss: 377854.9138
Epoch 35/35
1079/1079 [=====] - 2s 2ms/sample - loss: 374402.5451
lstm: [1512.103] 717.4, 900.9, 985.5, 1393.7, 1981.4, 1709.5, 2217.4

```



Saving figure Epochs35



## 5.4 Neurons/Layer and Number of Layers Experiment

### 5.4.1 100 neurons, 35 epochs, 2e/2d, dropout=0, regularization=none

```

[22]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 35, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(100, activation='relu', input_shape=(n_timesteps, ↪
    ↪n_features)))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(100, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(50, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size, ↪
    ↪verbose=verbose)
    return model

# evaluate model and get scores
score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("Neurons100Decoder2")
pyplot.show()

```

Train on 1079 samples

Epoch 1/35

```


1079/1079 [=====] - 2s 2ms/sample - loss:
1627782614.9842
Epoch 2/35
1079/1079 [=====] - 1s 1ms/sample - loss:
234017739.5737
Epoch 3/35
1079/1079 [=====] - 1s 1ms/sample - loss:
737185221.9018
Epoch 4/35
1079/1079 [=====] - 1s 1ms/sample - loss: 91449836.7238
Epoch 5/35
1079/1079 [=====] - 1s 1ms/sample - loss:
286834062.5839
Epoch 6/35
1079/1079 [=====] - 1s 1ms/sample - loss: 70397934.9379
Epoch 7/35
1079/1079 [=====] - 1s 1ms/sample - loss: 25179404.9861
Epoch 8/35
1079/1079 [=====] - 1s 1ms/sample - loss: 17516880.2954
Epoch 9/35
1079/1079 [=====] - 1s 1ms/sample - loss: 9827677.7071
Epoch 10/35
1079/1079 [=====] - 1s 1ms/sample - loss: 5726022.8811
Epoch 11/35
1079/1079 [=====] - 1s 1ms/sample - loss: 8396755.8184
Epoch 12/35
1079/1079 [=====] - 1s 1ms/sample - loss: 28105982.3772
Epoch 13/35
1079/1079 [=====] - 1s 1ms/sample - loss: 3956935.5904
Epoch 14/35
1079/1079 [=====] - 1s 1ms/sample - loss: 1955337.5589
Epoch 15/35
1079/1079 [=====] - 1s 1ms/sample - loss: 2033391.3274
Epoch 16/35
1079/1079 [=====] - 1s 1ms/sample - loss: 1056955.3782
Epoch 17/35
1079/1079 [=====] - 1s 1ms/sample - loss: 619958.1498
Epoch 18/35
1079/1079 [=====] - 1s 1ms/sample - loss: 495600.8310
Epoch 19/35
1079/1079 [=====] - 1s 1ms/sample - loss: 467774.8921
Epoch 20/35
1079/1079 [=====] - 1s 1ms/sample - loss: 490902.1154
Epoch 21/35
1079/1079 [=====] - 1s 1ms/sample - loss: 454644.7244
Epoch 22/35
1079/1079 [=====] - 1s 1ms/sample - loss: 484920.9901
Epoch 23/35

```


```

1079/1079 [=====] - 1s 1ms/sample - loss: 427613.8303
Epoch 24/35
1079/1079 [=====] - 1s 1ms/sample - loss: 422705.5261
Epoch 25/35
1079/1079 [=====] - 1s 1ms/sample - loss: 558554.77530s
- loss: 540
Epoch 26/35
1079/1079 [=====] - 1s 1ms/sample - loss: 636593.8871
Epoch 27/35
1079/1079 [=====] - 1s 1ms/sample - loss: 492904.7366
Epoch 28/35
1079/1079 [=====] - 1s 1ms/sample - loss: 440761.2313
Epoch 29/35
1079/1079 [=====] - 1s 1ms/sample - loss: 626609.7626
Epoch 30/35
1079/1079 [=====] - 1s 1ms/sample - loss: 9251150.6617
Epoch 31/35
1079/1079 [=====] - 1s 1ms/sample - loss: 2250818.5990
Epoch 32/35
1079/1079 [=====] - 1s 1ms/sample - loss: 950889.5856
Epoch 33/35
1079/1079 [=====] - 1s 1ms/sample - loss: 451186.3139
Epoch 34/35
1079/1079 [=====] - 1s 1ms/sample - loss: 627423.7670
Epoch 35/35
1079/1079 [=====] - 1s 1ms/sample - loss: 558761.8910
lstm: [760.547] 715.6, 669.7, 605.7, 812.6, 727.9, 767.7, 970.6

```



Saving figure Neurons100Decoder2



#### 5.4.2 400 neurons, 35 epochs, 2e/3d, dropout=0, regularization=none

```
[23]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 35, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(400, activation='relu', input_shape=(n_timesteps, ↪
    ↪n_features)))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(400, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(200, activation='relu')))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size, ↪
    ↪verbose=verbose)
```

```

    return model

# evaluate model and get scores
score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("Neurons400Decoder3")
pyplot.show()

```

Train on 1079 samples

Epoch 1/35

1079/1079 [=====] - 9s 9ms/sample - loss: 350945652.1816

Epoch 2/35

1079/1079 [=====] - 8s 7ms/sample - loss: 152675130.0019

Epoch 3/35

1079/1079 [=====] - 8s 7ms/sample - loss: 51594299.7461

Epoch 4/35

1079/1079 [=====] - 8s 7ms/sample - loss: 7607225.7257

Epoch 5/35

1079/1079 [=====] - 8s 7ms/sample - loss: 3793197.6766

Epoch 6/35

1079/1079 [=====] - 8s 7ms/sample - loss: 6173842.3939

Epoch 7/35

1079/1079 [=====] - 8s 7ms/sample - loss: 3272474.3698

Epoch 8/35

1079/1079 [=====] - 8s 7ms/sample - loss: 2533232.3935

Epoch 9/35

1079/1079 [=====] - 8s 7ms/sample - loss: 893101.7550

Epoch 10/35


1079/1079 [=====] - 8s 7ms/sample - loss: 526664.6292

Epoch 11/35




1079/1079 [=====] - 9s 8ms/sample - loss: 480065.7954  
 Epoch 12/35  
 1079/1079 [=====] - 9s 8ms/sample - loss: 445161.6005  
 Epoch 13/35  
 1079/1079 [=====] - 9s 8ms/sample - loss: 644560.8407  
 Epoch 14/35  
 1079/1079 [=====] - 9s 8ms/sample - loss: 613007.6233  
 Epoch 15/35  
 1079/1079 [=====] - 9s 9ms/sample - loss: 1442087.6704  
 Epoch 16/35  
 1079/1079 [=====] - 9s 8ms/sample - loss: 1043323.6124  
 Epoch 17/35  
 1079/1079 [=====] - 8s 8ms/sample - loss: 712946.9689  
 Epoch 18/35  
 1079/1079 [=====] - 8s 8ms/sample - loss: 408659.9425  
 Epoch 19/35  
 1079/1079 [=====] - 9s 8ms/sample - loss: 407012.4942  
 Epoch 20/35  
 1079/1079 [=====] - 9s 8ms/sample - loss: 396053.8325  
 Epoch 21/35  
 1079/1079 [=====] - 8s 8ms/sample - loss: 396733.0584  
 Epoch 22/35  
 1079/1079 [=====] - 8s 8ms/sample - loss: 420808.4667  
 Epoch 23/35  
 1079/1079 [=====] - 8s 8ms/sample - loss: 400491.6283  
 Epoch 24/35  
 1079/1079 [=====] - 9s 8ms/sample - loss: 405867.2980  
 Epoch 25/35  
 1079/1079 [=====] - 9s 8ms/sample - loss: 414049.7230  
 Epoch 26/35  
 1079/1079 [=====] - 8s 7ms/sample - loss: 380021.3067  
 Epoch 27/35  
 1079/1079 [=====] - 8s 7ms/sample - loss: 408374.1941  
 Epoch 28/35  
 1079/1079 [=====] - 8s 7ms/sample - loss: 378242.5232  
 Epoch 29/35  
 1079/1079 [=====] - 8s 7ms/sample - loss: 378265.6125  
 Epoch 30/35  
 1079/1079 [=====] - 8s 7ms/sample - loss: 427112.9139  
 Epoch 31/35  
 1079/1079 [=====] - 8s 7ms/sample - loss: 440628.6283  
 Epoch 32/35  
 1079/1079 [=====] - 8s 7ms/sample - loss: 371036.2166  
 Epoch 33/35  
 1079/1079 [=====] - 8s 7ms/sample - loss: 388836.2415  
 Epoch 34/35  
 1079/1079 [=====] - 8s 7ms/sample - loss: 394667.3049  
 Epoch 35/35

1079/1079 [=====] - 8s 8ms/sample - loss: 383025.5851  
lstm: [642.431] 619.8, 497.2, 459.6, 658.6, 613.2, 637.6, 911.1



Saving figure Neurons400Decoder3



### 5.4.3 200 neurons, 35 epochs, 2e/3d, dropout=0, regularization=none

```
[24]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 35, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps,
    ↪n_features)))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(200, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(50, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size,
    ↪verbose=verbose)
    return model

# evaluate model and get scores
score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("Neurons200Decoder3")
pyplot.show()
```

Train on 1079 samples

Epoch 1/35

1079/1079 [=====] - 3s 3ms/sample - loss:

707894145.3939

Epoch 2/35

1079/1079 [=====] - 2s 2ms/sample - loss: 94293044.7192

Epoch 3/35

1079/1079 [=====] - 2s 2ms/sample - loss: 26987750.6988

Epoch 4/35

1079/1079 [=====] - 2s 2ms/sample - loss: 24410645.7813

Epoch 5/35

1079/1079 [=====] - 2s 2ms/sample - loss: 10024179.9796

Epoch 6/35

1079/1079 [=====] - 2s 2ms/sample - loss: 7423889.5473

Epoch 7/35

1079/1079 [=====] - 2s 2ms/sample - loss: 3553172.1066

Epoch 8/35

1079/1079 [=====] - 2s 2ms/sample - loss: 2838227.3230

Epoch 9/35

1079/1079 [=====] - 2s 2ms/sample - loss: 5861982.9568

Epoch 10/35

1079/1079 [=====] - 2s 2ms/sample - loss: 685636.8706

Epoch 11/35

1079/1079 [=====] - 2s 2ms/sample - loss: 409219.8582

Epoch 12/35

1079/1079 [=====] - 2s 2ms/sample - loss: 417782.0028

Epoch 13/35

1079/1079 [=====] - 2s 2ms/sample - loss: 486914.6483

Epoch 14/35

1079/1079 [=====] - 2s 2ms/sample - loss: 2084411.6520

Epoch 15/35

1079/1079 [=====] - 2s 2ms/sample - loss: 476126.8503

Epoch 16/35

1079/1079 [=====] - 2s 2ms/sample - loss: 705277.6773

Epoch 17/35

1079/1079 [=====] - 2s 2ms/sample - loss: 623076.9990

Epoch 18/35

1079/1079 [=====] - 2s 2ms/sample - loss: 412000.7300

Epoch 19/35

1079/1079 [=====] - 2s 2ms/sample - loss: 344465.7478

Epoch 20/35

1079/1079 [=====] - 2s 2ms/sample - loss: 348548.9464

Epoch 21/35

1079/1079 [=====] - 2s 2ms/sample - loss: 401714.8907


Epoch 22/35

1079/1079 [=====] - 2s 2ms/sample - loss: 515632.4787


Epoch 23/35

1079/1079 [=====] - 2s 2ms/sample - loss: 563763.6064

Epoch 24/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 375826.2583  
 Epoch 25/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 346237.0240  
 Epoch 26/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 12443716.0957  
 Epoch 27/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 624149.2333  
 Epoch 28/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 748606.8209  
 Epoch 29/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 538422.1625  
 Epoch 30/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 450235.5433  
 Epoch 31/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 378405.7483  
 Epoch 32/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 391837.2130  
 Epoch 33/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 402442.3420  
 Epoch 34/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 436096.0206  
 Epoch 35/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 449057.7634  
 lstm: [1054.489] 745.4, 930.9, 1128.9, 1241.0, 1508.3, 579.9, 967.3



Saving figure Neurons200Decoder3



#### 5.4.4 200 neurons, 35 epochs, 2e/4d, dropout=0, regularization=none

```
[25]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 35, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps, ↪
    ↪n_features)))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(200, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(50, activation='relu')))
    model.add(TimeDistributed(Dense(10, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size, ↪
    ↪verbose=verbose)
```

```

    return model

# evaluate model and get scores
score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("Neurons200Decoder4")
pyplot.show()

```

Train on 1079 samples

Epoch 1/35

1079/1079 [=====] - 4s 3ms/sample - loss: 86268628.0946

Epoch 2/35

1079/1079 [=====] - 2s 2ms/sample - loss: 1458124.6680

Epoch 3/35

1079/1079 [=====] - 2s 2ms/sample - loss: 757906.1076

Epoch 4/35

1079/1079 [=====] - 2s 2ms/sample - loss: 597573.4613

Epoch 5/35

1079/1079 [=====] - 2s 2ms/sample - loss: 412961.4434

Epoch 6/35

1079/1079 [=====] - 2s 2ms/sample - loss: 434960.3990

Epoch 7/35

1079/1079 [=====] - 2s 2ms/sample - loss: 761531.4112

Epoch 8/35

1079/1079 [=====] - 2s 2ms/sample - loss: 1168756.2475

0s - loss: 109

Epoch 9/35

1079/1079 [=====] - 2s 2ms/sample - loss: 742192.9487

Epoch 10/35

1079/1079 [=====] - 2s 2ms/sample - loss: 726680.1602


Epoch 11/35

1079/1079 [=====] - 2s 2ms/sample - loss: 1817271.8937


Epoch 12/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2908921.9826  
Epoch 13/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2773927.2023  
Epoch 14/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2773658.5945  
Epoch 15/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2773397.3281  
Epoch 16/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2773143.1295  
Epoch 17/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2772894.5308  
Epoch 18/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2772651.4071  
Epoch 19/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2772411.7539  
Epoch 20/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2772175.1138  
Epoch 21/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2771941.1937  
Epoch 22/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2771709.9036  
0s - loss:  
Epoch 23/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2771480.5313  
Epoch 24/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2771253.1545  
Epoch 25/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2771027.4509  
Epoch 26/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2770803.1082  
Epoch 27/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2770579.8610  
Epoch 28/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2770357.8356  
Epoch 29/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2770136.9824  
Epoch 30/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2769917.0968  
Epoch 31/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2769697.5053  
Epoch 32/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2769479.0630  
Epoch 33/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2769261.3406  
Epoch 34/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2769044.4516  
Epoch 35/35



1079/1079 [=====] - 2s 2ms/sample - loss: 2768828.0462  
lstm: [1567.987] 1612.4, 1527.5, 1517.0, 1597.7, 1495.5, 1505.4, 1708.9



Saving figure Neurons200Decoder4



## 5.5 Dropout Experiment

### 5.5.1 200 neurons, 35 epochs, 2e/3d, dropout=.2, regularization=None

```
[26]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 35, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps,
    ↪n_features), dropout=0.2))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(200, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(50, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size,
    ↪verbose=verbose)
    return model

# evaluate model and get scores
score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("Dropout2")
```

```
pyplot.show()
```

Train on 1079 samples

Epoch 1/35

1079/1079 [=====] - 3s 3ms/sample - loss: 336494132.9082

Epoch 2/35

1079/1079 [=====] - 2s 2ms/sample - loss: 146973501.5829

Epoch 3/35

1079/1079 [=====] - 2s 2ms/sample - loss: 72940269.0371

Epoch 4/35

1079/1079 [=====] - 2s 2ms/sample - loss: 41429232.5283

Epoch 5/35

1079/1079 [=====] - 2s 2ms/sample - loss: 17697709.1047

Epoch 6/35

1079/1079 [=====] - 2s 2ms/sample - loss: 13939509.3068

Epoch 7/35

1079/1079 [=====] - 2s 2ms/sample - loss: 19310171.3559

Epoch 8/35

1079/1079 [=====] - 2s 2ms/sample - loss: 5419528.3957

Epoch 9/35

1079/1079 [=====] - 2s 2ms/sample - loss: 3721178.4893

Epoch 10/35

1079/1079 [=====] - 2s 2ms/sample - loss: 7574442.5058

Epoch 11/35

1079/1079 [=====] - 2s 2ms/sample - loss: 6635565.2845

Epoch 12/35

1079/1079 [=====] - 2s 2ms/sample - loss: 40219594.2261

Epoch 13/35

1079/1079 [=====] - 2s 2ms/sample - loss: 12040307.5412

Epoch 14/35

1079/1079 [=====] - 2s 2ms/sample - loss: 4101439.4092

Epoch 15/35

1079/1079 [=====] - 2s 2ms/sample - loss: 3360666.3647

Epoch 16/35

1079/1079 [=====] - 2s 2ms/sample - loss: 3201969.8781

1s - loss: 406 - ETA: 1s -

Epoch 17/35

1079/1079 [=====] - 2s 2ms/sample - loss: 3042377.2397

Epoch 18/35

1079/1079 [=====] - 2s 2ms/sample - loss: 8270607.1353

Epoch 19/35

1079/1079 [=====] - 2s 2ms/sample - loss: 4542311.1657

Epoch 20/35


1079/1079 [=====] - 2s 2ms/sample - loss: 9778778.3424

Epoch 21/35


```

1079/1079 [=====] - 2s 2ms/sample - loss: 4722361.8577
0s - loss: 50024
Epoch 22/35
1079/1079 [=====] - 2s 2ms/sample - loss: 1790007.3007
Epoch 23/35
1079/1079 [=====] - 2s 2ms/sample - loss: 1380768.2443
Epoch 24/35
1079/1079 [=====] - 2s 2ms/sample - loss: 1135490.7287
Epoch 25/35
1079/1079 [=====] - 2s 2ms/sample - loss: 46158421.5630
Epoch 26/35
1079/1079 [=====] - 2s 2ms/sample - loss: 31392845.1622
Epoch 27/35
1079/1079 [=====] - 2s 2ms/sample - loss: 10857038.1872
Epoch 28/35
1079/1079 [=====] - 2s 2ms/sample - loss: 2704954.6543
Epoch 29/35
1079/1079 [=====] - 2s 2ms/sample - loss: 2646432.5132
Epoch 30/35
1079/1079 [=====] - 2s 2ms/sample - loss: 1683931.5591
Epoch 31/35
1079/1079 [=====] - 2s 2ms/sample - loss: 989042.3570:
0s - loss: 997631.4
Epoch 32/35
1079/1079 [=====] - 2s 2ms/sample - loss: 816400.0897
Epoch 33/35
1079/1079 [=====] - 2s 2ms/sample - loss: 537434.5203
Epoch 34/35
1079/1079 [=====] - 2s 2ms/sample - loss: 544228.1121
Epoch 35/35
1079/1079 [=====] - 2s 2ms/sample - loss: 634798.5151
lstm: [826.815] 847.0, 815.0, 789.9, 905.3, 761.5, 705.5, 939.5

```



Saving figure Dropout2



### 5.5.2 200 neurons, 50 epochs, 2e/3d, dropout=.2, regularization=none

```
[27]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 50, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps, ↪
    ↪n_features), dropout=0.2))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(200, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(50, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size, ↪
    ↪verbose=verbose)
    return model

# evaluate model and get scores
score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("Dropout2Epochs50")
pyplot.show()
```

Train on 1079 samples

Epoch 1/50

1079/1079 [=====] - 3s 3ms/sample - loss:

1995901374.8434

Epoch 2/50

1079/1079 [=====] - 2s 2ms/sample - loss:

267867988.7303

Epoch 3/50

1079/1079 [=====] - 2s 2ms/sample - loss:

739245055.0806

Epoch 4/50

1079/1079 [=====] - 2s 2ms/sample - loss:

245106326.9509

Epoch 5/50

1079/1079 [=====] - 2s 2ms/sample - loss:

40631119.46250s - loss: 426

Epoch 6/50

1079/1079 [=====] - 2s 2ms/sample - loss: 30002704.7868

Epoch 7/50

1079/1079 [=====] - 2s 2ms/sample - loss: 65485354.1307

Epoch 8/50

1079/1079 [=====] - 2s 2ms/sample - loss: 10732740.1274

Epoch 9/50

1079/1079 [=====] - 2s 2ms/sample - loss: 5415792.1353

Epoch 10/50

1079/1079 [=====] - 2s 2ms/sample - loss: 14014867.1886

Epoch 11/50

1079/1079 [=====] - 2s 2ms/sample - loss: 7663816.9911

Epoch 12/50

1079/1079 [=====] - 2s 2ms/sample - loss: 718103.0678

Epoch 13/50

1079/1079 [=====] - 2s 2ms/sample - loss: 633914.0090

Epoch 14/50

1079/1079 [=====] - 2s 2ms/sample - loss: 786763.46211s

- loss: 764

Epoch 15/50

1079/1079 [=====] - 2s 2ms/sample - loss: 1411347.8962

Epoch 16/50

1079/1079 [=====] - 2s 2ms/sample - loss: 852403.8428

Epoch 17/50

1079/1079 [=====] - 2s 2ms/sample - loss: 595046.5699

Epoch 18/50

1079/1079 [=====] - 2s 2ms/sample - loss: 530112.9220

Epoch 19/50

1079/1079 [=====] - 2s 2ms/sample - loss: 498598.7812

Epoch 20/50


1079/1079 [=====] - 2s 2ms/sample - loss: 621022.8597

Epoch 21/50


1079/1079 [=====] - 2s 2ms/sample - loss: 651593.8879  
Epoch 22/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 698155.9774  
Epoch 23/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 2221597.4993  
Epoch 24/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 2093521.1084  
Epoch 25/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1465749.5477  
Epoch 26/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 2187845.2632  
Epoch 27/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1838889.1931  
Epoch 28/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 3462679.2121  
Epoch 29/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 3609606.9662  
Epoch 30/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 2864259.1513  
Epoch 31/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1617354.8583  
Epoch 32/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1122545.9412  
Epoch 33/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 932504.4332  
Epoch 34/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 912273.7871  
Epoch 35/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1529346.4818  
Epoch 36/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1131417.6081  
Epoch 37/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 2795513.0807  
Epoch 38/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1214362.6463  
Epoch 39/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1397560.5142  
Epoch 40/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1075900.3716  
Epoch 41/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1469175.4835  
Epoch 42/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 1996135.2405  
Epoch 43/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 2492841.9985  
Epoch 44/50  
1079/1079 [=====] - 2s 2ms/sample - loss: 5372248.1664  
Epoch 45/50



```
1079/1079 [=====] - 2s 2ms/sample - loss: 1432699.0987
Epoch 46/50
1079/1079 [=====] - 2s 2ms/sample - loss: 1177905.0623
Epoch 47/50
1079/1079 [=====] - 2s 2ms/sample - loss: 1082110.5191
Epoch 48/50
1079/1079 [=====] - 2s 2ms/sample - loss: 3216760.5673
Epoch 49/50
1079/1079 [=====] - 2s 2ms/sample - loss: 919339.2563
Epoch 50/50
1079/1079 [=====] - 2s 2ms/sample - loss: 810732.4088
lstm: [776.465] 807.7, 597.8, 640.5, 847.3, 783.9, 785.5, 922.4
```



Saving figure Dropout2Epochs50



### 5.5.3 200 neurons, 35 epochs, 2e/3d, dropout=.5, regularization=none

```
[28]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 35, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps, ↪
    ↪n_features), dropout=0.5))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(200, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(50, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size, ↪
    ↪verbose=verbose)
    return model

# evaluate model and get scores
```

```

score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("Dropout5")
pyplot.show()

```

Train on 1079 samples

Epoch 1/35

1079/1079 [=====] - 4s 3ms/sample - loss: 786056076.7618

Epoch 2/35

1079/1079 [=====] - 2s 2ms/sample - loss: 445256146.4467

Epoch 3/35

1079/1079 [=====] - 2s 2ms/sample - loss: 104598025.8703

Epoch 4/35

1079/1079 [=====] - 2s 2ms/sample - loss: 90512104.1372

Epoch 5/35

1079/1079 [=====] - 2s 2ms/sample - loss: 41953099.5292

Epoch 6/35

1079/1079 [=====] - 2s 2ms/sample - loss: 65098452.8193

Epoch 7/35

1079/1079 [=====] - 2s 2ms/sample - loss: 37980035.1066

Epoch 8/35

1079/1079 [=====] - 2s 2ms/sample - loss: 72341442.6006

Epoch 9/35

1079/1079 [=====] - 2s 2ms/sample - loss: 137443572.9787

Epoch 10/35


1079/1079 [=====] - 2s 2ms/sample - loss: 48438639.7201

Epoch 11/35


1079/1079 [=====] - 2s 2ms/sample - loss: 7620573.9618

Epoch 12/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 3142702.1475  
Epoch 13/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 4733105.2690  
Epoch 14/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 4667252.1209  
Epoch 15/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 4625052.4690  
Epoch 16/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 4935781.7224  
Epoch 17/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2862752.0770  
Epoch 18/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 9817330.2451  
Epoch 19/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 9431127.3960  
Epoch 20/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 11743353.2634  
Epoch 21/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 13702035.9949  
Epoch 22/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 3370478.3777  
Epoch 23/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2473446.7785  
Epoch 24/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 4256765.7303  
Epoch 25/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2020270.9269  
Epoch 26/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2992985.6419  
Epoch 27/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2686251.5642  
Epoch 28/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2307016.8797  
Epoch 29/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2188158.6637  
Epoch 30/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 1716343.0346  
Epoch 31/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 3568900.0199  
Epoch 32/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2018180.2609  
1s -  
Epoch 33/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 1343906.4578  
Epoch 34/35  
1079/1079 [=====] - 2s 2ms/sample - loss: 2519370.7944  
Epoch 35/35

1079/1079 [=====] - 2s 2ms/sample - loss: 1611114.2024  
lstm: [1123.419] 944.6, 903.0, 836.8, 1011.5, 1103.0, 1282.3, 1594.5



Saving figure Dropout5



## 5.6 Regularization Experiment

### 5.6.1 200 neurons, 35 epochs, 2e/3d, dropout=0, kernel regularization=l1(0.01)

```
[29]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 35, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps,
    ↪n_features), kernel_regularizer='l1'))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(200, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(50, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size,
    ↪verbose=verbose)
    return model

# evaluate model and get scores
score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("RegularizationL1")
```

```
pyplot.show()
```

Train on 1079 samples

Epoch 1/35

1079/1079 [=====] - 3s 3ms/sample - loss:

499188318.2280 0s - loss: 593886859.555 - ETA: 0s - loss: 57581303

Epoch 2/35

1079/1079 [=====] - 2s 2ms/sample - loss: 86749841.0528

Epoch 3/35

1079/1079 [=====] - 2s 2ms/sample - loss: 69373701.0408

Epoch 4/35

1079/1079 [=====] - 2s 2ms/sample - loss: 19016579.9147

Epoch 5/35

1079/1079 [=====] - 2s 2ms/sample - loss: 13456029.5894

Epoch 6/35

1079/1079 [=====] - 2s 2ms/sample - loss: 17706760.7488

Epoch 7/35

1079/1079 [=====] - 2s 2ms/sample - loss: 23702944.7859

Epoch 8/35

1079/1079 [=====] - 2s 2ms/sample - loss: 8017827.2034

Epoch 9/35

1079/1079 [=====] - 2s 2ms/sample - loss: 3669562.3377

Epoch 10/35

1079/1079 [=====] - 2s 2ms/sample - loss: 3840063.0491

Epoch 11/35

1079/1079 [=====] - 2s 2ms/sample - loss: 3439732.7878

1s

Epoch 12/35

1079/1079 [=====] - 2s 2ms/sample - loss: 2791639.8415

Epoch 13/35

1079/1079 [=====] - 2s 2ms/sample - loss: 5442931.2244

Epoch 14/35

1079/1079 [=====] - 2s 2ms/sample - loss: 4604769.0424

Epoch 15/35

1079/1079 [=====] - 2s 2ms/sample - loss: 719007.3466

Epoch 16/35

1079/1079 [=====] - 2s 2ms/sample - loss: 340195.4933

Epoch 17/35

1079/1079 [=====] - 2s 2ms/sample - loss: 324290.4778

Epoch 18/35

1079/1079 [=====] - 2s 2ms/sample - loss: 398026.2073

Epoch 19/35

1079/1079 [=====] - 2s 2ms/sample - loss: 320154.4936

Epoch 20/35


1079/1079 [=====] - 2s 2ms/sample - loss: 323730.6725

Epoch 21/35


1079/1079 [=====] - 2s 2ms/sample - loss: 803102.8077

```
Epoch 22/35
1079/1079 [=====] - 2s 2ms/sample - loss: 8004000.3368
Epoch 23/35
1079/1079 [=====] - 2s 2ms/sample - loss: 4128955.9886
Epoch 24/35
1079/1079 [=====] - 2s 2ms/sample - loss: 4458794.1562
Epoch 25/35
1079/1079 [=====] - 2s 2ms/sample - loss: 5039320.4989
Epoch 26/35
1079/1079 [=====] - 2s 2ms/sample - loss: 432919.6908
Epoch 27/35
1079/1079 [=====] - 2s 2ms/sample - loss: 646364.1550
Epoch 28/35
1079/1079 [=====] - 2s 2ms/sample - loss: 1697684.7815
Epoch 29/35
1079/1079 [=====] - 2s 2ms/sample - loss: 1435750.7033
Epoch 30/35
1079/1079 [=====] - 2s 2ms/sample - loss: 1511628.5215
Epoch 31/35
1079/1079 [=====] - 2s 2ms/sample - loss: 1006345.2500
Epoch 32/35
1079/1079 [=====] - 2s 2ms/sample - loss: 753314.7709
Epoch 33/35
1079/1079 [=====] - 2s 2ms/sample - loss: 590069.1573
Epoch 34/35
1079/1079 [=====] - 2s 2ms/sample - loss: 596061.1783
Epoch 35/35
1079/1079 [=====] - 2s 2ms/sample - loss: 777020.5209
lstm: [773.021] 796.4, 699.1, 617.5, 792.2, 662.1, 828.2, 962.5
```





Saving figure RegularizationL1



### 5.6.2 200 neurons, 35 epochs, 2e/3d, dropout=0, kernel regularization=l2(0.01)

```
[30]: # train the model
def build_model(train, n_input):
    # prepare data
    train_x, train_y = to_supervised(train, n_input)
    # define parameters
    verbose, epochs, batch_size = 1, 35, 16
    n_timesteps, n_features, n_outputs = train_x.shape[1], train_x.
    ↪shape[2], train_y.shape[1]
    # reshape output into [samples, timesteps, features]
    train_y = train_y.reshape((train_y.shape[0], train_y.shape[1], 1))
    # define model
    model = Sequential()
    model.add(LSTM(200, activation='relu', input_shape=(n_timesteps,
    ↪n_features), kernel_regularizer='l2'))
    model.add(RepeatVector(n_outputs))
    model.add(LSTM(200, activation='relu', return_sequences=True))
    model.add(TimeDistributed(Dense(100, activation='relu')))
    model.add(TimeDistributed(Dense(50, activation='relu')))
    model.add(TimeDistributed(Dense(1)))
    model.compile(loss='mse', optimizer='adam')
    # fit network
    model.fit(train_x, train_y, epochs=epochs, batch_size=batch_size,
    ↪verbose=verbose)
    return model

# evaluate model and get scores
score, scores, predictions = evaluate_model(train, test, n_input)
# summarize scores
summarize_scores('lstm', score, scores)
predictions_df = pd.DataFrame.from_records(predictions)
predicted_values = predictions_df.mean(axis=0)

# plot scores
days = ['sun', 'mon', 'tue', 'wed', 'thr', 'fri', 'sat']
pyplot.plot(days, scores, marker='o', label='lstm')
pyplot.show()

# plot actual vs predicted
pyplot.figure(figsize=(10,5))
pyplot.plot(days, predicted_values, label='predicted')
pyplot.plot(days, test_values, label='actual')
pyplot.legend()
pyplot.title("Average Power Consumption by Day of the Week (46 weeks)")
save_fig("RegularizationL2")
pyplot.show()
```

Train on 1079 samples

Epoch 1/35

1079/1079 [=====] - 4s 3ms/sample - loss:

390137304.3855

Epoch 2/35

1079/1079 [=====] - 2s 2ms/sample - loss: 83415887.4217

Epoch 3/35

1079/1079 [=====] - 2s 2ms/sample - loss: 51482507.8109

Epoch 4/35

1079/1079 [=====] - 2s 2ms/sample - loss: 48847138.7044

Epoch 5/35

1079/1079 [=====] - 2s 2ms/sample - loss: 7832066.4022

Epoch 6/35

1079/1079 [=====] - 2s 2ms/sample - loss: 1686235.4233

Epoch 7/35

1079/1079 [=====] - 2s 2ms/sample - loss: 3250003.2694

Epoch 8/35

1079/1079 [=====] - 2s 2ms/sample - loss: 412364.3508

Epoch 9/35

1079/1079 [=====] - 2s 2ms/sample - loss: 411216.2292

Epoch 10/35

1079/1079 [=====] - 2s 2ms/sample - loss: 632167.5381

Epoch 11/35

1079/1079 [=====] - 2s 2ms/sample - loss: 937123.5877

Epoch 12/35

1079/1079 [=====] - 2s 2ms/sample - loss: 573836.1207

Epoch 13/35

1079/1079 [=====] - 2s 2ms/sample - loss: 435778.1305

Epoch 14/35

1079/1079 [=====] - 2s 2ms/sample - loss: 456151.6732

Epoch 15/35

1079/1079 [=====] - 2s 2ms/sample - loss: 420654.9846

Epoch 16/35

1079/1079 [=====] - 2s 2ms/sample - loss: 484272.6482

Epoch 17/35

1079/1079 [=====] - 2s 2ms/sample - loss: 465339.2096

Epoch 18/35

1079/1079 [=====] - 2s 2ms/sample - loss: 413044.2812

Epoch 19/35

1079/1079 [=====] - 2s 2ms/sample - loss: 415085.1830

Epoch 20/35

1079/1079 [=====] - 2s 2ms/sample - loss: 526066.1262

Epoch 21/35

1079/1079 [=====] - 2s 2ms/sample - loss: 411028.7176


Epoch 22/35

1079/1079 [=====] - 2s 2ms/sample - loss: 451853.9906


Epoch 23/35

1079/1079 [=====] - 2s 2ms/sample - loss: 408929.1913

Epoch 24/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 464919.3236  
 Epoch 25/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 508247.9908  
 Epoch 26/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 537332.0120  
 Epoch 27/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 440161.6779  
 Epoch 28/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 461588.4599  
 Epoch 29/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 461890.7542  
 Epoch 30/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 442068.8170  
 Epoch 31/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 458724.3975  
 Epoch 32/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 6629323.8445  
 Epoch 33/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 1207304.5990  
 Epoch 34/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 11369418.3313  
 Epoch 35/35  
 1079/1079 [=====] - 2s 2ms/sample - loss: 3200509.9328  
 0s - loss:  
 lstm: [1160.710] 910.4, 767.5, 920.7, 1204.8, 1180.2, 1335.0, 1593.3



Saving figure RegularizationL2



[ ]: